

# CLASIFICACIÓN MULTIPLE - POPULARIDAD

## Integrantes de equipo:

- Julio César Choquehuayta Quenta 2018-119025
- Carlos Enrique Yufra Loza 2019-119051
- Pablo Moisés Aro Galindo 2019-119034
- Carlos Manuel Azañero Otoyá 2015-119026
- Jose Angel Castro Caceres 2017-119054

## DESCRIPCION DEL PROBLEMA

La industria de la musica es bastante popular en la actualidad, como lo que escucha las personajes cambia con el tiempo, ya sea por tendencias o moda, mucho de lo que determina su popularidad suele ser subjetivo o circunstancial. Sin embargo tomando en cuenta datos tecnicos o variables presentes en la musica, se tendra como objetivo el determinar si es popular o no.

Este proyecto se enfocara en un problema de clasificacion multiple, el cual hara uso de una dataset de Spotify de pistas en un rango de 125 géneros diferentes. Cada pista tiene algunas funciones de audio asociadas. Los datos están en formato CSV, que es tabular y se puede cargar rápidamente.

Link del Dataset: <https://www.kaggle.com/datasets/maharshipandya/-spotify-tracks-dataset?resource=download>

## ANALISIS EXPLORATORIO

Primero, nosotros necesitamos tener la dataset descargada y almacenada en una carpeta en drive para su proximo uso, ademas de importar las librerias necesarias.

```
1 import numpy as np
2 import pandas as pd
3 import tensorflow as tf
4 from tensorflow import keras
5 from tensorflow.keras import layers,regularizers
6 from sklearn.model_selection import train_test_split
7 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
8 from keras.models import Sequential
9 from keras.layers import Dense
10 from sklearn.metrics import accuracy_score, confusion_matrix
11 import matplotlib.pyplot as plt
```

Nosotros cargamos la dataset usando pandas read\_csv, junto con el link de donde esta nuestra data en drive

```
1 # Obtenemos el dataset de las Canciones de Spotify y lo almacenamos en un DataFrame
2 df = pd.read_csv("https://drive.google.com/uc?id=1gAkvuBkYAEKCdf9CYqePni8TacqkwWt3")
3 df
```

| Unnamed: 0 |        | track_id               | artists                | album_name   | track_name                 | popularity | duration_ms |
|------------|--------|------------------------|------------------------|--|----------------------------|------------|-------------|
| 0          | 0      | 5SuOikwiRyPMVoIQDJUgSV | Gen Hoshino            | Comedy   | Comedy                     | 73         | 230666      |
| 1          | 1      | 4qPNDBW1i3p13qLCt0Ki3A | Ben Woodward           | Ghost (Acoustic)                                       | Ghost - Acoustic           | 55         | 149610      |
| 2          | 2      | 1iJBSr7s7jYXzM8EGcbK5b | Ingrid Michaelson;ZAYN | To Begin Again   | To Begin Again             | 57         | 210826      |
| 3          | 3      | 6lfqx3CG4xtTiEg7opyCyx | Kina Grannis           | Crazy Rich Asians (Original Motion Picture Soundtrack) | Can't Help Falling In Love | 71         | 201933      |
| 4          | 4      | 5vjLSffimilP26QG5WcN2K | Chord Overstreet       | Hold On  | Hold On                    | 82         | 198853      |
| ...        | ...    | ...                    | ...                    | ...  | ...                        | ...        | ...         |
| 443335     | 443335 | 6COTF7ED1A1V5P1K37     | Billie Eilish          | #mindfulness - Soft Rain                               | Sleep Mv                   | 64         | 201000      |

PRE-PROCESAMIENTO DE LOS DATOS

```
##### CARGA DE DATOS #####
##### CARGA DE DATOS #####

1 # Convertimos la columna 'popularity' en 'popularity_class'(0:baja, 1:media, 2:alta)
2 df['popularity_class'] = 0
3 df.loc[df['popularity'] > 30, 'popularity_class'] = 1
4 df.loc[df['popularity'] > 70, 'popularity_class'] = 2

1 # Extraemos una muestra de 100 000 registros aleatorios
2 df_sample = df.sample(100000)
3
4 # Filtramos las entradas y la salida que usaremos para el modelo
5 columns = ['danceability','energy','loudness','speechiness','acousticness','instrumentalness','valence','tempo','popularity_class']
6 df_columns = df_sample[columns]
7 df_columns.nunique() # visualizamos la cantidad de valores únicos de las columnas

danceability      1151
energy            2041
loudness          18807
speechiness       1478
acousticness      4995
instrumentalness  5326
valence           1777
tempo            42471
popularity_class    3
dtype: int64

1 # Visualizamos algunos registros filtrados por el valor de 'popularity_class'
2 df_columns.query('popularity_class == 1')
```

|        | danceability | energy | loudness | speechiness | acousticness | instrumentalness | valence | tempo   | popularity_class |
|--------|--------------|--------|----------|-------------|--------------|------------------|---------|---------|------------------|
| 23775  | 0.731        | 0.988  | -5.174   | 0.1220      | 0.00251      | 0.799000         | 0.155   | 125.993 | 1                |
| 9365   | 0.661        | 0.529  | -5.330   | 0.0312      | 0.64800      | 0.000000         | 0.328   | 121.892 | 1                |
| 104241 | 0.304        | 0.328  | -10.311  | 0.0351      | 0.56000      | 0.000000         | 0.243   | 119.675 | 1                |
| 72140  | 0.582        | 0.928  | -4.659   | 0.0635      | 0.00876      | 0.000005         | 0.565   | 113.043 | 1                |
| 77744  | 0.504        | 0.743  | -6.529   | 0.1050      | 0.55500      | 0.000000         | 0.707   | 90.425  | 1                |
| ...    | ...          | ...    | ...      | ...         | ...          | ...              | ...     | ...     | ...              |
| 47429  | 0.521        | 0.721  | -7.699   | 0.0254      | 0.00348      | 0.011300         | 0.611   | 97.578  | 1                |
| 105246 | 0.736        | 0.289  | -11.398  | 0.0494      | 0.93000      | 0.839000         | 0.150   | 68.987  | 1                |
| 89948  | 0.664        | 0.893  | -3.627   | 0.1500      | 0.07480      | 0.000000         | 0.773   | 104.007 | 1                |
| 19813  | 0.666        | 0.252  | -17.614  | 0.0264      | 0.54800      | 0.022100         | 0.547   | 94.111  | 1                |
| 12978  | 0.439        | 0.446  | -8.113   | 0.0319      | 0.68100      | 0.000000         | 0.300   | 142.837 | 1                |

50287 rows × 9 columns

```
1 # Visualizamos algunas estadísticas de nuestros datos
2 df_columns.describe()

count    100000.000000    100000.000000    100000.000000    100000.000000    100000.000000    100000.000000    100000.000000
mean      0.566666      0.641036      -8.260964      0.084568      0.315636      0.156273      0.47408
std       0.173611      0.251811      5.034996      0.105598      0.332711      0.309847      0.25935
min       0.000000      0.000000     -49.531000      0.000000      0.000000      0.000000      0.00000
25%       0.456000      0.471000     -10.013000      0.035900      0.017000      0.000000      0.26000
50%       0.580000      0.684000      -7.004000      0.048900      0.170000      0.000042      0.46400
75%       0.694000      0.854000      -5.004000      0.084400      0.599000      0.049000      0.68300
max       0.985000      1.000000      4.532000      0.965000      0.996000      1.000000      0.99400

1 # Pasamos los datos del DF a un arreglo
2 data = df_columns.values
3
4 # Separamos las características de las etiquetas
5 x_data = data[:, :-1]
6 y_data = data[:, -1]
7
8 print("X:")
9 print(x_data)
10 print("y:")
11 print(y_data)

X:
[[ 5.23000e-01  9.12000e-01 -2.24300e+00 ...  0.00000e+00  8.24000e-01
  1.61906e+02]
 [ 5.17000e-01  6.09000e-01 -5.70200e+00 ...  0.00000e+00  2.39000e-01
  9.20930e+01]
 [ 7.31000e-01  9.88000e-01 -5.17400e+00 ...  7.99000e-01  1.55000e-01
  1.25993e+02]
 ...
 [ 6.66000e-01  2.52000e-01 -1.76140e+01 ...  2.21000e-02  5.47000e-01
  9.41110e+01]
 [ 4.39000e-01  4.46000e-01 -8.11300e+00 ...  0.00000e+00  3.00000e-01
  1.42837e+02]
 [ 4.20000e-01  6.46000e-01 -7.61500e+00 ...  1.53000e-02  8.77000e-01
  1.42209e+02]]
y:
[0. 0. 1. ... 1. 1. 0.]

1 # Verificamos las dimensiones de los arreglos
2 print("X:")
3 print(x_data.shape)
4 print("y:")
5 print(y_data.shape)

X:
(100000, 8)
y:
(100000,)
```

```
1 # Obtenemos la media y la desviación estándar de cada característica
2 x_mean = x_data.mean(axis = 0)
3 x_std = x_data.std(axis = 0)
4
5 # Normalizamos las características del modelo y convertimos a float32
6 x_data = (x_data - x_mean) / x_std
7 x_data= x_data.astype(np.float32)
8
9 # Verificamos los datos normalizados
10 x_data

array([[ -0.25151557,  1.0760666 ,  1.1952331 , ..., -0.5043599 ,
         1.3491517 ,  1.3268516 ],
       [ -0.2860758 , -0.1272223 ,  0.5082381 , ..., -0.5043599 ,
        -0.9064179 , -1.000805 ],
       [ 0.946572 ,  1.3778816 ,  0.6131046 , ...,  2.0743487 ,
        -1.2302946 ,  0.12946522],
       ...,
       [ 0.5721696 , -1.5449587 , -1.8576149 , ..., -0.43303394,
```

https://colab.research.google.com/drive/17ewjzUQnimB9paZlwkfAn\_yHt112hui0#scrollTo=lgchCWBqyR5d&printMode=true

```

    0.28112987, -0.93352234],
[-0.7353586 , -0.77453613,  0.02938719, ..., -0.5043599 ,
-0.6712218 ,  0.6910662 ],
[-0.84479934,  0.01971398,  0.12829542, ..., -0.45498037,
 1.5535026 ,  0.67012787]], dtype=float32)

1 # Separamos aleatoriamente los datos de entrenamiento(80%) y de prueba(20%)
2 X_train, X_test, y_train, y_test = train_test_split(x_data, y_data, test_size=0.2, random_state=20)

1 # Definimos el modelo
2 model = Sequential()
3 model.add(Dense(512, activation='relu', input_shape=(8,)))
4 model.add(layers.Dropout(0.125))
5 model.add(Dense(512, activation='relu'))
6 # model.add(layers.Dropout(0.125))
7 model.add(Dense(512, activation='relu'))
8 # model.add(layers.Dropout(0.125))
9 model.add(Dense(512*1, activation='relu'))
10 model.add(Dense(3, activation='softmax'))
11
12 # Compilamos el modelo
13 model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics='accuracy')
14
15 # Entrenamos el modelo
16 history = model.fit(X_train, y_train, validation_split=0.2, epochs=100, batch_size=1024)

Epoch 1/100
63/63 [=====] - 2s 9ms/step - loss: 0.8104 - accuracy: 0.5655 - val_loss: 0.7958 - val_accuracy: 0.5860
Epoch 2/100
63/63 [=====] - 0s 6ms/step - loss: 0.7830 - accuracy: 0.5914 - val_loss: 0.7855 - val_accuracy: 0.5954
Epoch 3/100
63/63 [=====] - 0s 5ms/step - loss: 0.7759 - accuracy: 0.6007 - val_loss: 0.7784 - val_accuracy: 0.5989
Epoch 4/100
63/63 [=====] - 0s 5ms/step - loss: 0.7705 - accuracy: 0.6031 - val_loss: 0.7803 - val_accuracy: 0.5960
Epoch 5/100
63/63 [=====] - 0s 6ms/step - loss: 0.7685 - accuracy: 0.6081 - val_loss: 0.7751 - val_accuracy: 0.6044
Epoch 6/100
63/63 [=====] - 0s 6ms/step - loss: 0.7625 - accuracy: 0.6135 - val_loss: 0.7716 - val_accuracy: 0.6060
Epoch 7/100
63/63 [=====] - 0s 5ms/step - loss: 0.7587 - accuracy: 0.6162 - val_loss: 0.7689 - val_accuracy: 0.6089
Epoch 8/100
63/63 [=====] - 0s 5ms/step - loss: 0.7568 - accuracy: 0.6169 - val_loss: 0.7668 - val_accuracy: 0.6133
Epoch 9/100
63/63 [=====] - 1s 9ms/step - loss: 0.7531 - accuracy: 0.6203 - val_loss: 0.7661 - val_accuracy: 0.6090
Epoch 10/100
63/63 [=====] - 1s 8ms/step - loss: 0.7504 - accuracy: 0.6219 - val_loss: 0.7618 - val_accuracy: 0.6150
Epoch 11/100
63/63 [=====] - 0s 8ms/step - loss: 0.7460 - accuracy: 0.6283 - val_loss: 0.7630 - val_accuracy: 0.6190
Epoch 12/100
63/63 [=====] - 0s 8ms/step - loss: 0.7415 - accuracy: 0.6300 - val_loss: 0.7605 - val_accuracy: 0.6235
Epoch 13/100
63/63 [=====] - 1s 8ms/step - loss: 0.7377 - accuracy: 0.6345 - val_loss: 0.7582 - val_accuracy: 0.6286
Epoch 14/100
63/63 [=====] - 1s 8ms/step - loss: 0.7325 - accuracy: 0.6370 - val_loss: 0.7550 - val_accuracy: 0.6227
Epoch 15/100
63/63 [=====] - 1s 10ms/step - loss: 0.7278 - accuracy: 0.6415 - val_loss: 0.7532 - val_accuracy: 0.6258
Epoch 16/100
63/63 [=====] - 0s 7ms/step - loss: 0.7240 - accuracy: 0.6441 - val_loss: 0.7588 - val_accuracy: 0.6236
Epoch 17/100
63/63 [=====] - 0s 6ms/step - loss: 0.7201 - accuracy: 0.6455 - val_loss: 0.7512 - val_accuracy: 0.6276
Epoch 18/100
63/63 [=====] - 0s 5ms/step - loss: 0.7114 - accuracy: 0.6528 - val_loss: 0.7503 - val_accuracy: 0.6311
Epoch 19/100
63/63 [=====] - 0s 5ms/step - loss: 0.7070 - accuracy: 0.6553 - val_loss: 0.7525 - val_accuracy: 0.6256
Epoch 20/100
63/63 [=====] - 0s 5ms/step - loss: 0.7013 - accuracy: 0.6581 - val_loss: 0.7531 - val_accuracy: 0.6244
Epoch 21/100
63/63 [=====] - 0s 5ms/step - loss: 0.6941 - accuracy: 0.6630 - val_loss: 0.7531 - val_accuracy: 0.6316
Epoch 22/100
63/63 [=====] - 0s 5ms/step - loss: 0.6871 - accuracy: 0.6706 - val_loss: 0.7441 - val_accuracy: 0.6349
Epoch 23/100
63/63 [=====] - 0s 5ms/step - loss: 0.6790 - accuracy: 0.6718 - val_loss: 0.7533 - val_accuracy: 0.6296
Epoch 24/100
63/63 [=====] - 0s 5ms/step - loss: 0.6727 - accuracy: 0.6744 - val_loss: 0.7554 - val_accuracy: 0.6343
Epoch 25/100
63/63 [=====] - 0s 5ms/step - loss: 0.6652 - accuracy: 0.6792 - val_loss: 0.7468 - val_accuracy: 0.6409
Epoch 26/100
63/63 [=====] - 0s 5ms/step - loss: 0.6577 - accuracy: 0.6846 - val_loss: 0.7552 - val_accuracy: 0.6340
Epoch 27/100
63/63 [=====] - 0s 5ms/step - loss: 0.6474 - accuracy: 0.6911 - val_loss: 0.7569 - val_accuracy: 0.6390
Epoch 28/100

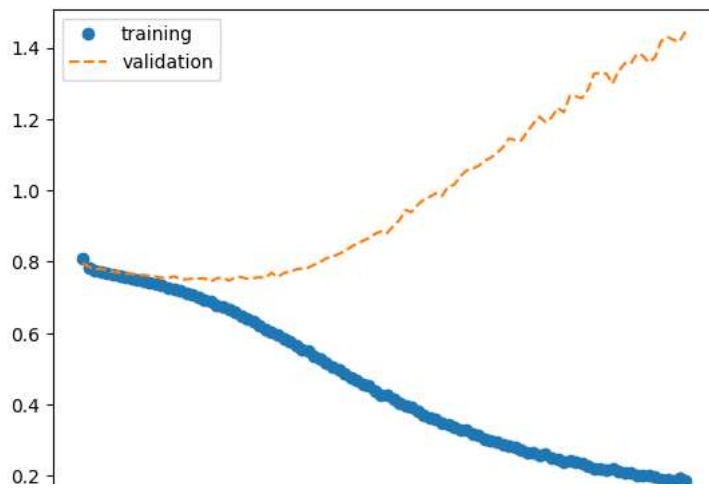
```

```
63/63 [=====] - 0s 5ms/step - loss: 0.6394 - accuracy: 0.6936 - val_loss: 0.7518 - val_accuracy: 0.6418
Epoch 29/100
63/63 [=====] - 0s 5ms/step - loss: 0.6319 - accuracy: 0.6986 - val_loss: 0.7540 - val_accuracy: 0.6424
```

```
1 # Guardamos los datos del entrenamiento
2 history_dict = history.history
3 history_dict
```

```
{'loss': [0.81043541431427,
0.783037543296814,
0.7758663892745972,
0.7705049514770508,
0.7685242891311646,
0.7624958157539368,
0.758703351020813,
0.7567902207374573,
0.7531402707099915,
0.7504097819328308,
0.746046781539917,
0.7415114045143127,
0.7376968264579773,
0.7325431108474731,
0.7278305292129517,
0.724018394947052,
0.7201322317123413,
0.7114274501800537,
0.7070450782775879,
0.7013481855392456,
0.6941251158714294,
0.6871050000190735,
0.679010272026062,
0.6727263331413269,
0.6652444005012512,
0.6577200889587402,
0.6474365592002869,
0.6393818855285645,
0.6318901181221008,
0.6209021210670471,
0.6115968823432922,
0.6022710800170898,
0.595801591873169,
0.58481764793396,
0.5761724710464478,
0.5635596513748169,
0.5544502139091492,
0.5483798980712891,
0.5335936546325684,
0.5265988707542419,
0.5144517421722412,
0.5052611827850342,
0.4970416724681854,
0.4858230650424957,
0.47665101289749146,
0.4662398397922516,
0.45778217911720276,
0.4506412744522095,
0.4391242563724518,
0.4277111291885376,
0.42490658164024353,
0.41359302401542664,
0.40362775325775146,
0.3968048691749573,
0.3917872905731201,
0.38151267170906067,
0.37148287892341614,
0.3641209900379181,
```


```
1 # Guardamos los valores de perdida en dos arreglos para entrenamiento y validacion
2 i=0
3 loss_values = history_dict["loss"][i:]
4 val_loss_values = history_dict["val_loss"][i:]
5
6 # Graficamos los valores de perdida
7 epoch = range(1, len(loss_values) + 1)
8 plt.plot(epoch, loss_values, 'o', label = 'training')
9 plt.plot(epoch, val_loss_values, '--', label = 'validation')
10 plt.legend()
11 plt.show()
```



```

1 # Guardamos los valores de precisión en dos arreglos para entrenamiento y validacion
2 i=0
3 loss_values = history_dict["accuracy"][i:]
4 val_loss_values = history_dict["val_accuracy"][i:]
5
6 # Graficamos los valores de precisión
7 epoch = range(1, len(loss_values) + 1)
8 plt.plot(epoch, loss_values, 'o', label = 'training')
9 plt.plot(epoch, val_loss_values, '--', label = 'validation')
10 plt.legend()
11 plt.show()

```



```

1 # evaluamos el modelo
2 model.evaluate(X_test, y_test)

625/625 [=====] - 2s 2ms/step - loss: 1.4669 - accuracy: 0.6852
[1.4668784141540527, 0.6851500272750854]

```

```

1 # Predecimos en base a los datos de prueba
2 y_predict_prob = model.predict(X_test)
3
4 # Convertimos las probabilidades en clases
5 y_predict = np.argmax(y_predict_prob, axis=1)

625/625 [=====] - 1s 2ms/step

```

```

1 # Verificamos algunos casos individuales
2 pos = 31
3 print("y: ", y_test[pos])
4 print("y predict:", y_predict[pos])

y: 1.0
y predict: 1

```

```
1 # Matriz de Confusion
2 print("Accuracy: {:.2f}%".format(accuracy_score(y_test, y_predict) * 100))
3 print("Confusion Matrix:")
4 print(confusion_matrix(y_test, y_predict))
```

```
Accuracy: 68.52%
Confusion Matrix:
[[6300 2657 134]
 [2887 7026 168]
 [ 205  246 377]]
```

